

Training cost by stage, mean $\pm$ std over seeds
(four datasets ran concurrently, so absolute values include contention)

